

GRAPH, COLUMN, AND IN MEMORY DATA BASES

Master Program in Computer Science
University of Calabria

Prof. F. Ricca

Graph Databases (1)

- DB designs disagree on various aspects of design, **but:**
 - Databases store information about “things” and their relationships!
 - Relationships have as natural counterpart: **Graphs**
- Relational Databases can easily model graphs, **but**
 - *RDBMS generally hits performance issues with large graphs!*
 - *SQL might be a little limiting (no full recursion!)*
- NoSQL are even worse at modeling graphs
 - Graph structure can be stored in a single document or object, **but**
 - *relationships between objects are not inherently supported (no joins!)*
- The need for a different DBMS
 - **The Graph database system**

RDBMS Patterns for Graphs

- The relational model can easily represent the data that is contained in a graph model, but
 - SQL lacks the syntax to perform graph traversal
 - Each level of traversal adds significantly to query response time

```
1 SELECT p2.personname, m1.movieName
2 FROM people p1
3 JOIN actors a1 ON (p1.personid = a1.personid)
4 JOIN movies m1 ON (a1.movieid = m1.movieid)
5 JOIN actors a2 ON (a2.movieid = m1.movieid)
6 JOIN people p2 ON (p2.personid = a2.personid)
7 WHERE p1.personname = 'Keanu Reeves';
8
```

Figure 5-3. SQL to perform first-level graph traversal

Graph Databases (2)

- A Graph is given by
 - A set of Nodes (or vertices) **N** and
 - A set of edges (or arcs) **E** modeling associations between two nodes.
 - **N** and **E** can have associated properties (labelled graphs)
- There is mathematical notation for performing operations
 - Add or remove nodes and edges or find adjacent nodes, etc.
- "Queries" end up in performing a *graph traversal*
 - *i.e.*, walking through the graph to explore the graph

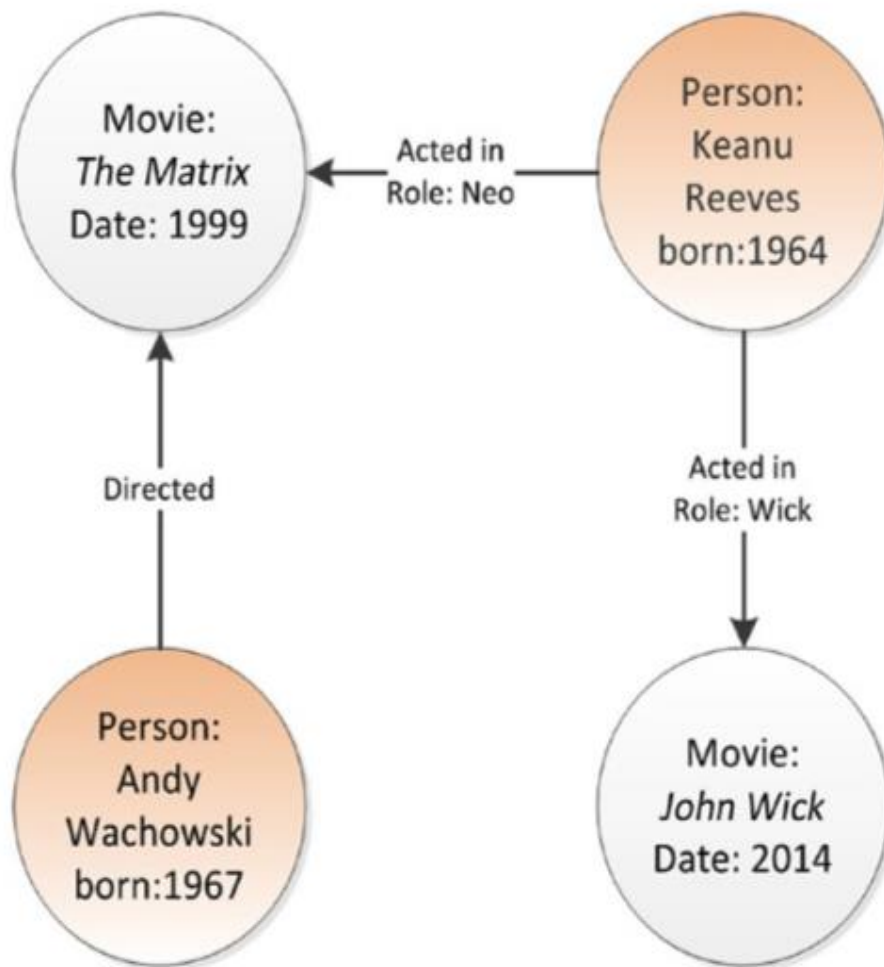


Figure 5-1. Simple graph with four vertices (nodes) and three edges (relationship)

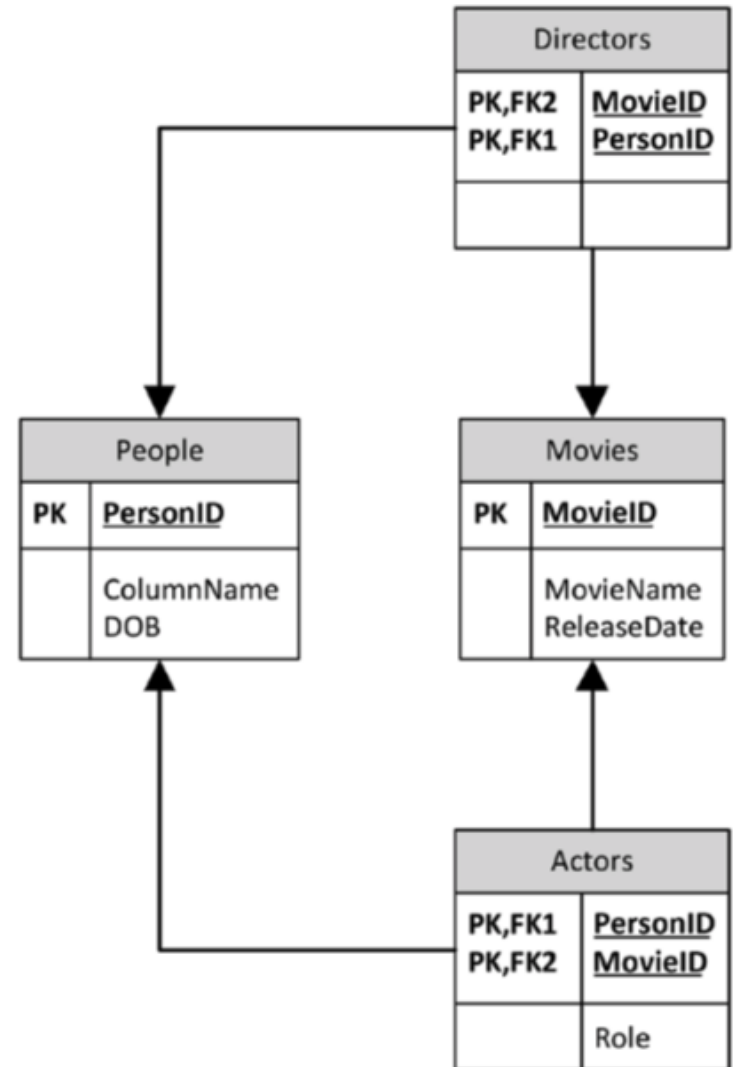


Figure 5-2. Relational schema to represent our sample graph data

Property Graphs and Neo4j

- *Property Graph* model
 - Associates both nodes and relationships with attributes
- *Neo4j*
 - Probably the most widely adopted graph database
 - *Cypher*: a declarative graph query language
 - Better than SPARQL to specify graph traversals

```
neo4j-sh (?)$ MATCH (kenau:Person {name:"Keanu Reeves"})  
> RETURN kenau;
```

- *Gremlin*: procedural graph database query language

Graph Database Efficient Structures

- *Index-free adjacency*
 - Enabling efficient real-time graph processing
 - Move through the graph structure without index lookups
 - Each node knows the physical location of all adjacent nodes
 - No need to use indexes to efficiently navigate the graph
- Not easy to distribute
 - The overhead of routing the traversal across multiple machines eliminates the advantages of the graph database model
 - Inter-server communication is far more time consuming than local access
 - Neo4j do not currently support a distributed deployment

Graph compute engines

- Graph processing and other data models work across massive distributed datasets
 - **Apache Giraph**
 - Designed to run over Hadoop data using MapReduce
 - **GraphX**
 - Part of the *Berkeley Data Analytic Stack (BDAS)*
 - Uses Spark as the foundation for graph processing
 - **Titan**
 - A scalable graph database optimized for storing and querying graphs distributed across a multi-machine cluster
 - Can be layered on top of Big Data storage engines, including Hbase and Cassandra

RDF

- The *Resource Description Framework (RDF)*
 - A web standard developed in the late 1990s
 - Information expressed in *triples*
 - entity: attribute :value
 - Intended for creating a formal database of resources
- RDF specification included XML syntax
- Native RDF databases are called *triplestores*
 - AllegroGraph, Ontotext GraphDB, StarDog, Virtuoso, and Oracle Spatial
- The *SPARQL* query language (a SQL-like language)

```
SELECT ?object
FROM <http://dbpedia.org>
WHERE { <http://dbpedia.org/resource/Edgar_F._Codd>
        <http://purl.org/dc/terms/subject> ?object }
```

Column Databases

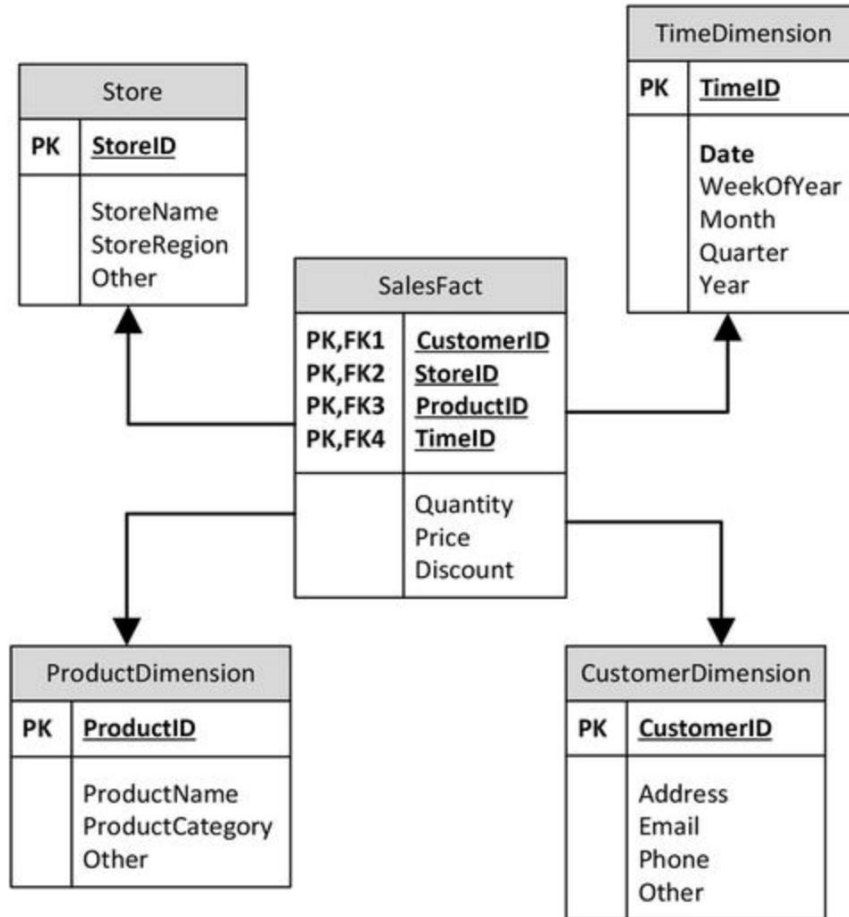
Those of us raised in Western cultures have been conditioned to think of data as arranged in rows
[Guy Harrison]

- The first databases addressed OLTP processing
 - CRUD operations were the most time-critical ones
 - Row-based storage (record-oriented) provided good performance
- Analytic and decision support applications
 - Data warehousing and analytic workloads
 - Long batch processing and aggregations
 - *OLAP (Online Analytic Processing)*
 - Coined by Edgar Codd to differentiate from OLTP systems

OLTP vs OLAP

- OLTP system
 - ACID transactions and CRUD operations
 - Service-level response times
 - Row-oriented physical organization ideal (write-intensive)
- OLAP system
 - IO intensive aggregate queries
 - *New non-normal-form schemas: Star & Snowflake*
 - Column-oriented physical organization ideal (read-intensive)
- Can we support both?
 - Analytic and decision support applications with interactive response times
 - C-Store (<http://db.lcs.mit.edu/projects/cstore/vldb.pdf>)

Star schema



- A schema for data warehouses:
 - Aggregate queries should execute quickly
- Central large *fact* tables are associated with
- Numerous smaller *dimension* tables
- Not a normalized relational model of data...

Figure 6-1. Star schema

Star schema and scalability

- Almost all data warehouses adopted some variation on the star schema paradigm
- Almost all DBMS adopted indexing and SQL optimizations to accelerate queries against star schemas
- However, *processing in data warehouses remained severely CPU and IO intensive*

The columnar Alternative

- Data for columns is grouped together on disk
- Clear advantage in aggregate queries, but:
- Retrieving a single row involves assembling the row from each of the column stores for that table

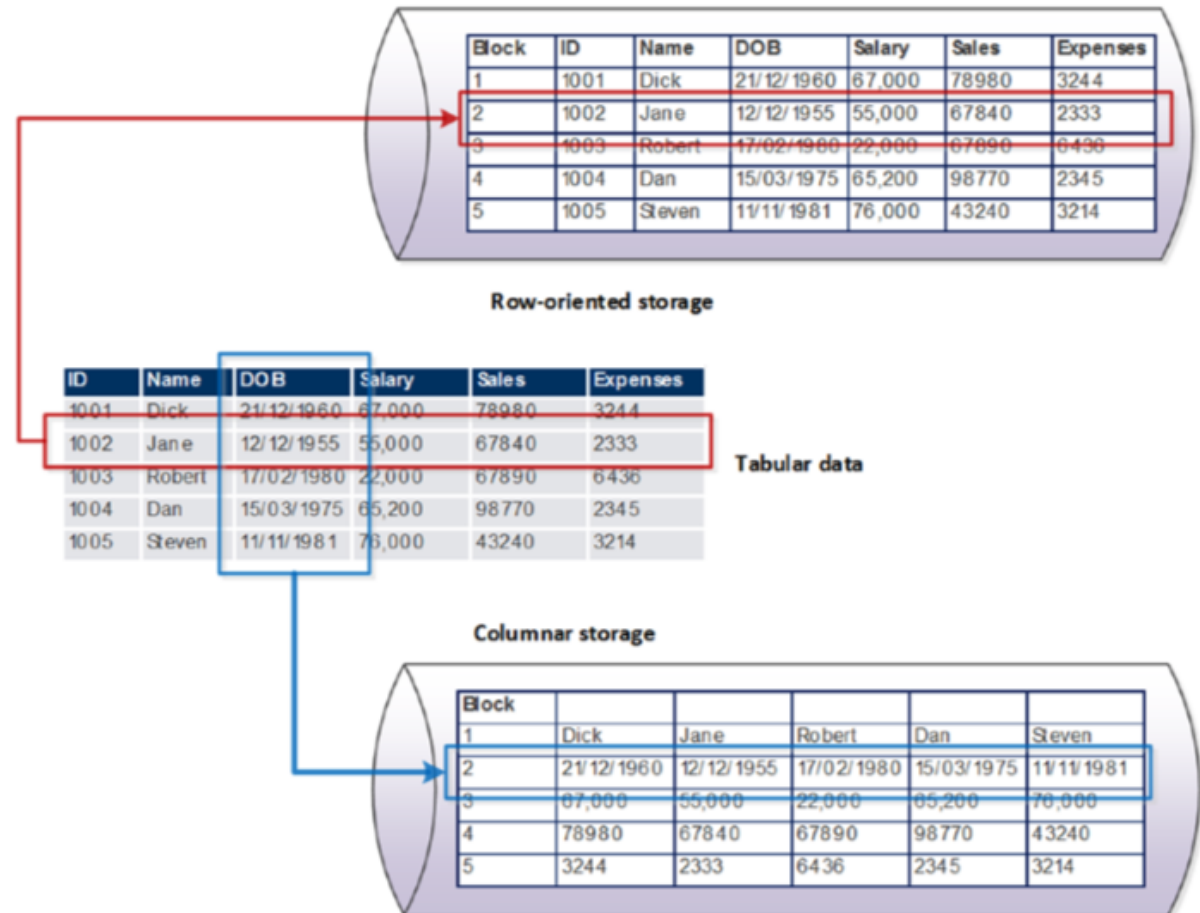


Figure 6-2. Comparison of columnar and row-oriented storage

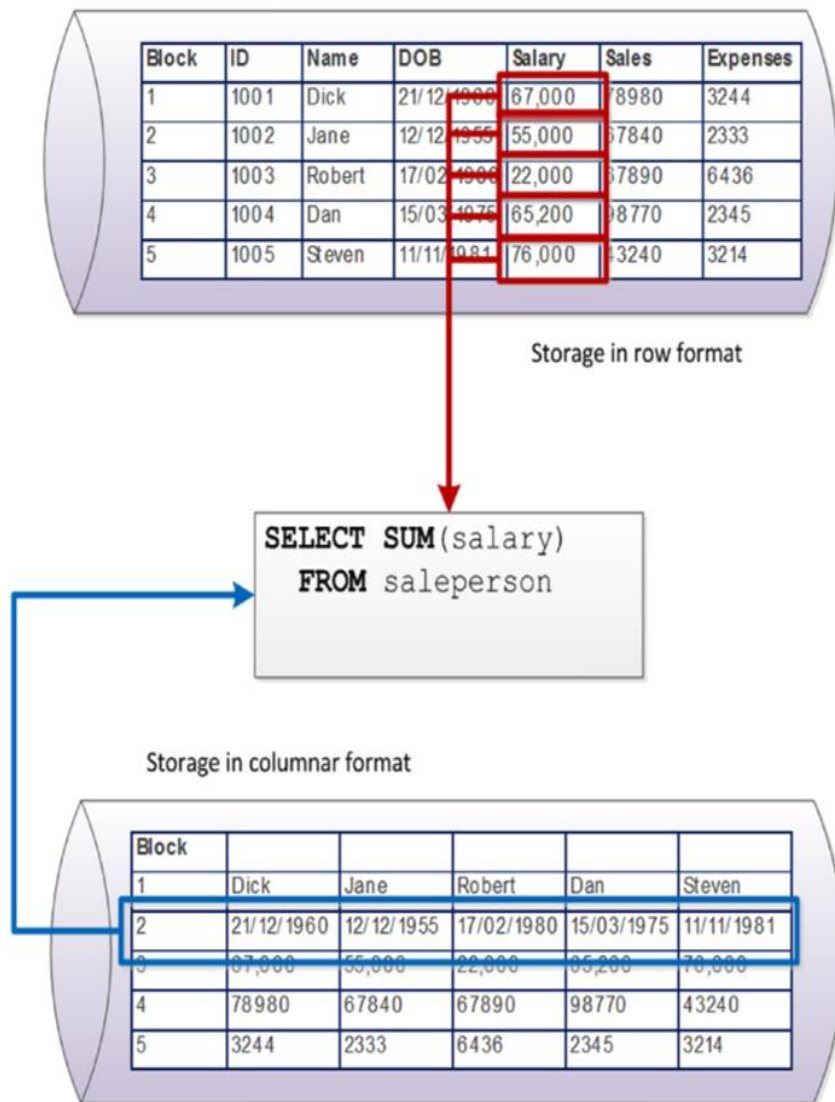


Figure 6-3. Aggregate operations in columnar stores require fewer IOs

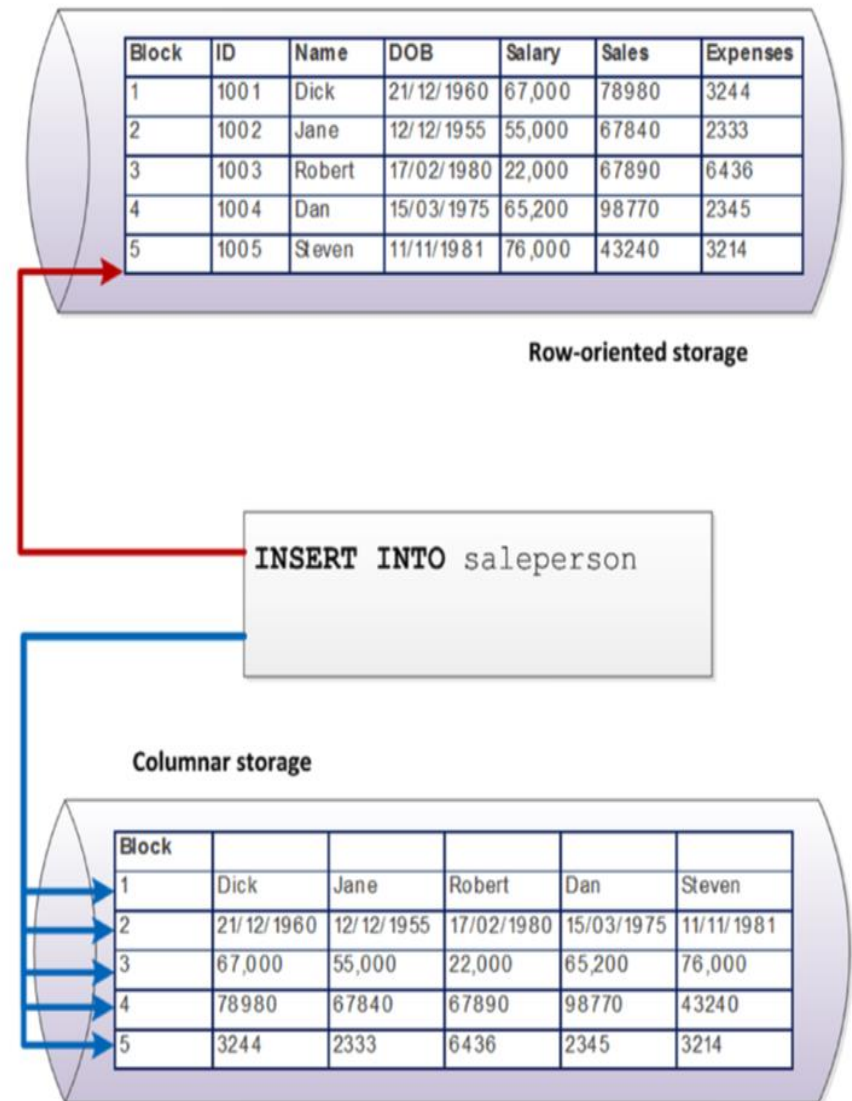


Figure 6-4. Insert overhead for a column store

Historical perspective (1)

- In the mid-1990s, Sybase acquired Expressway
 - The first significant commercial column-oriented database
 - Expressway became Sybase IQ (“Intelligent Query”)
 - *Sybase IQ failed to dominate the data warehousing market*
- In 2005, Mike Stonebraker and colleagues present C-store
 - A novel *column-oriented* DBMS
 - Faster than row-based DBMS for data warehousing

Historical perspective (2)

- Stonebraker founded *Vertica*
 - Acquired by HP in 2011
- Column-based systems entered the market
 - InfoBright, VectorWise, and MonetDB.
 - Oracle Exadata, Microsoft SQL Server, and SAP HANA
- Column-based systems
 - Significantly different from traditional RDBMS
 - Still based on SQL -> “NewSQL”

C-Store storage (1)

- Redundant storage of elements
 - Table split in overlapping *projections* in different orders
- Projections
 - Store combinations of columns together on disk
 - Combinations of columns that are frequently accessed together
 - Possibly sorted w.r.t. a different attribute
 - Redundant, but very efficient
- A set of projections covers the entire table
 - Additional projections are created to support specific queries

Region	Customer	Product	Sales
A	G	C	789
B	C	C	743
D	F	D	675
C	C	A	23
A	R	B	654

Logical Table

Table appears to user in relational normal form

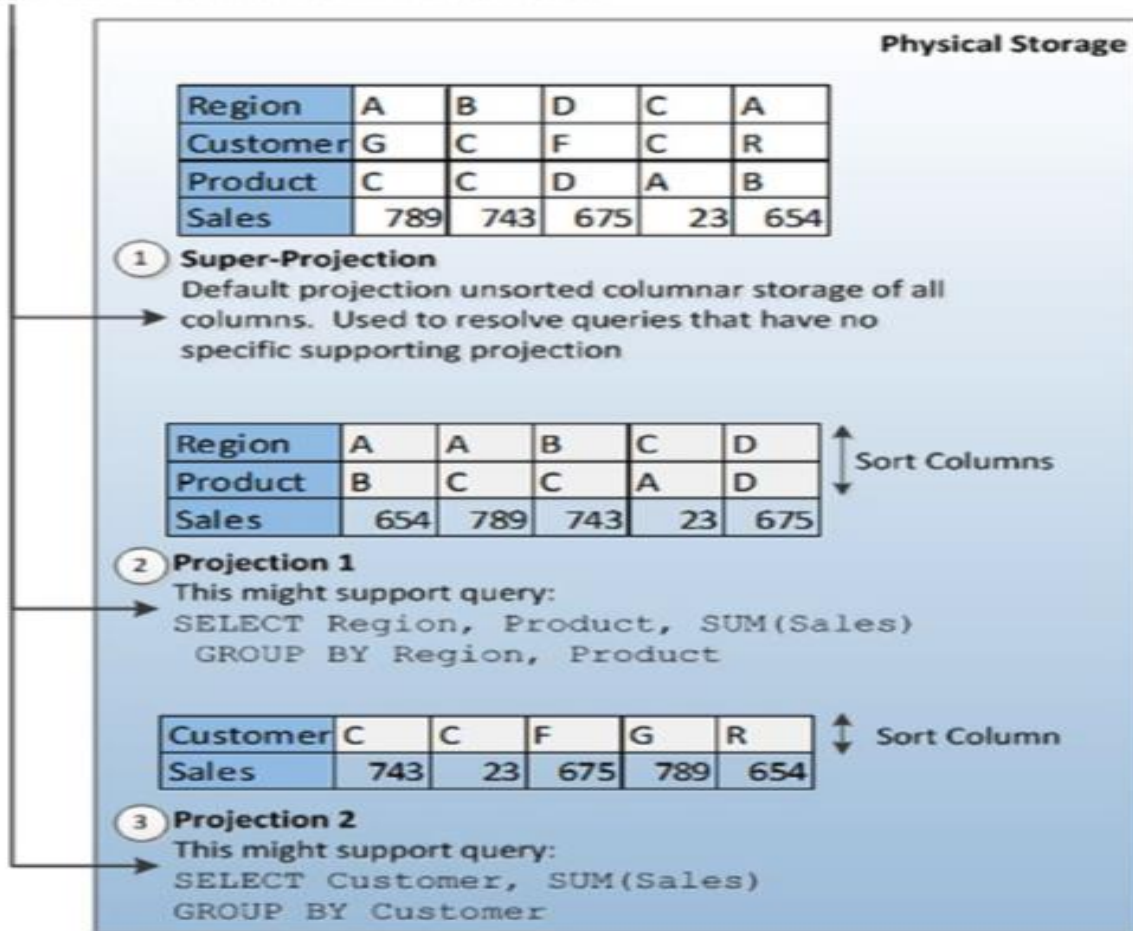
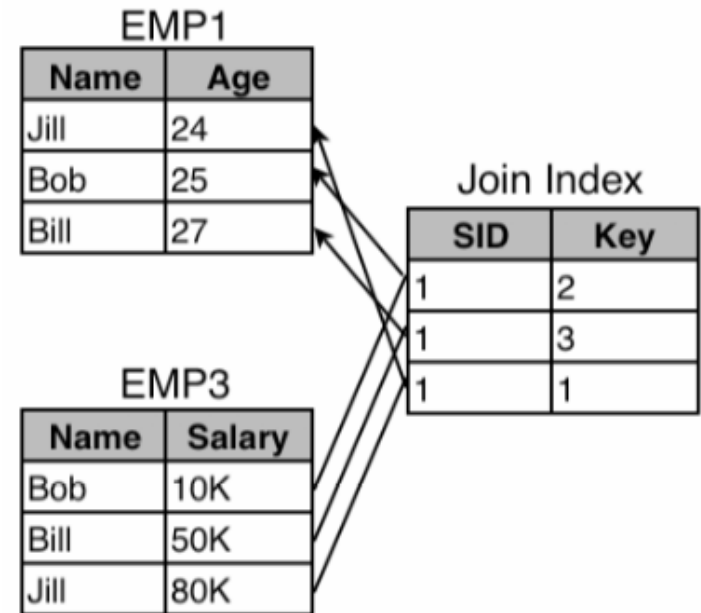


Figure 6-6. Columnar database table with three projections

C-Store storage (2)

- Data partitioning and distribution
 - Horizontally partitioned into 1 or more *segments*
 - Segments associate data values of every column with a *storage key*
 - *Join indices* reconstruct all the records in a table



C-Store storage (3)

- Columnar Compression
 - Usually try to work on localized subsets of the data
 - Data is stored in sorted order
 - Few “combinations of values” high compression is possible!
- Heavily compressed columns
 - Trade CPU time for I/O bandwidth!
 - CPU overhead of compression lower on isolated blocks of data (columns are stored together on disk)

C-Store optimizer and transactions

- **A hybrid architecture**
 - Frequent insert and update and query performance
 - Column-oriented optimizer and executor
- High availability and improved performance
 - K-safety using a sufficient number of overlapping projections
- The use of *snapshot isolation*, 2PC for writes
 - Timestamps for reads consistent in a given “snapshot”

Columns Stores hybrid architecture (1)

- *Writeable (Delta) Store*
 - Optimized for frequent writes
 - Like an in-memory row-oriented db
 - Periodically merged with the main columnar-oriented store by "*Tuple Mover*"
- *Read-Optimized Store*
 - *Main Column store*

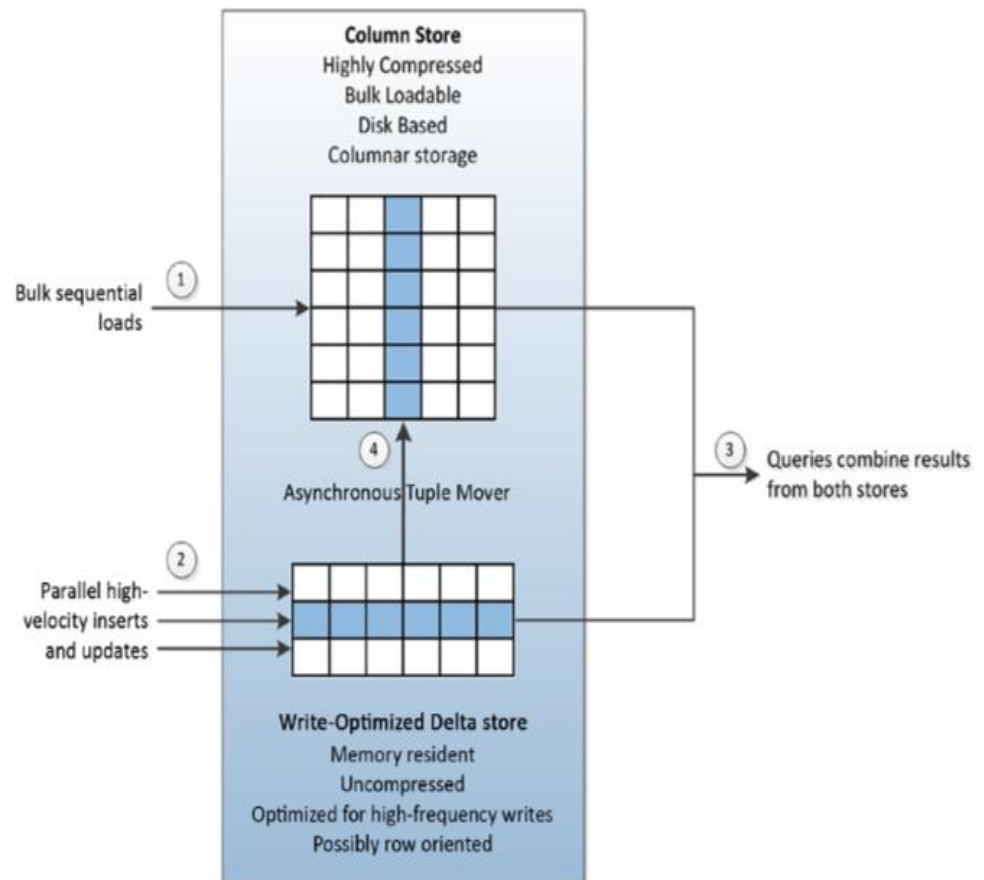


Figure 6-5. Write optimization in column store databases

Columns Stores hybrid architecture (2)

- Nightly ETL jobs
 - Column store
- Incremental inserts and updates
 - Write-optimized store
- Queries may need to read from both!

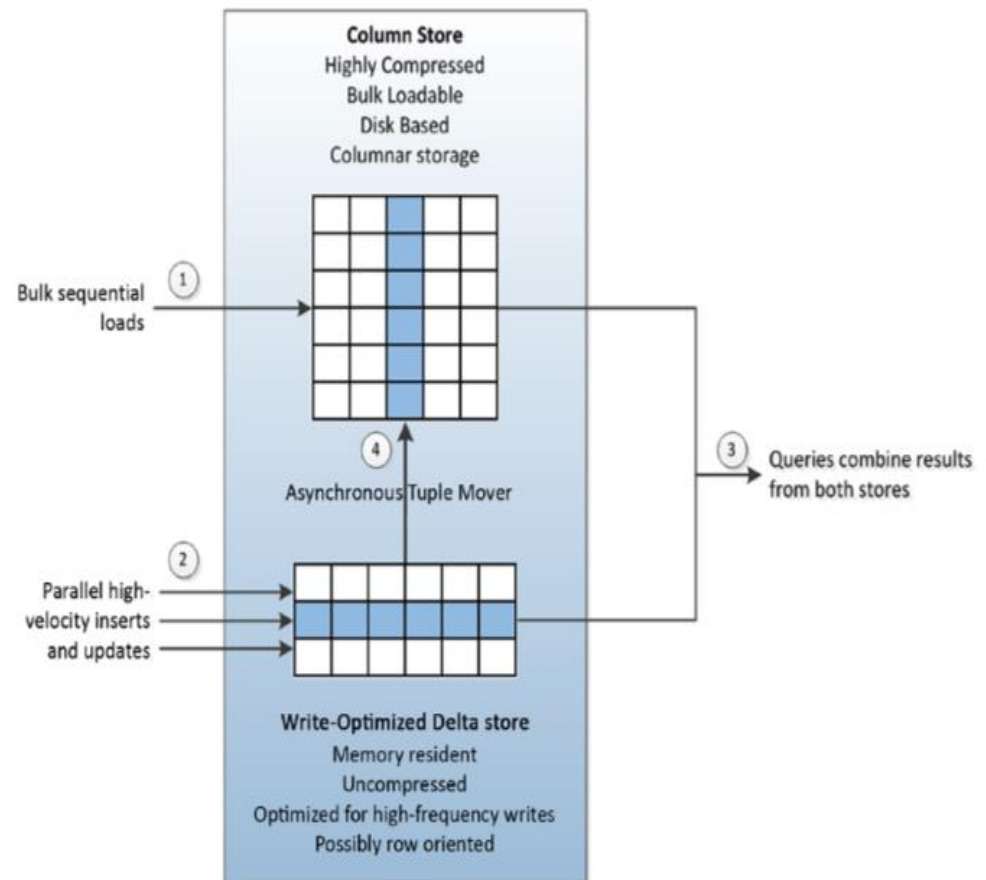


Figure 6-5. Write optimization in column store databases

Disks and Databases Architectures

- The magnetic disk device has been a pervasive presence within digital computing since the 1950s
 - Moore's law does not apply to the mechanical aspects of disk performance
 - But alternative technologies brought significant improvements!
- Solid state provide tremendously lower IO latencies
 - Performance of SSD is on orders of magnitude superior
 - Read operations require only a single-page IO
 - Writing a page requires an erase and overwrite of a complete block
 - *SSD is significantly slower in writes than in reads*




Figure 7-3. Trends in SSD and HDD storage costs (note logarithmic scale)

SSD and Databases

- Traditional relational databases perform relatively poorly on
 - Critical IO has been isolated to sequential write
 - The worst possible workload for SSD
- Some nonrelational systems
 - Avoid updating existing blocks and perform larger batched write operations (friendlier to solid state disk performance)
- *Aerospike – SSD-oriented*
 - NoSQL database
 - A log-structured file system
 - Updates are physically implemented by appending the new value to the file and marking the original data as invalid
 - Main memory to store indexes, data always on flash
 - “Avoid IO at all costs” approach unnecessary




Figure 7-4. Trends for memory cost and capacity for the last 40 years (note logarithmic scale)

Main memory DB

- **Cache-less architecture**

- Traditional disk-based databases cache data in main memory to minimize disk IO
- There is no point caching in memory what is already stored in memory!

- **Alternative persistence model**

- Data in memory disappears when the power is turned off
- Alternative mechanism for ensuring no data loss
 - Replicating data to other members of a cluster
 - Writing complete database images to disk files
 - Writing out transaction/operation records to a *transaction log* or *journal*

TimesTen (1995 – oracle 2005)

- SQL-based relational model
- Data is memory resident
- Periodic snapshots of memory to disk
- Disk-based transaction log following a transaction commit

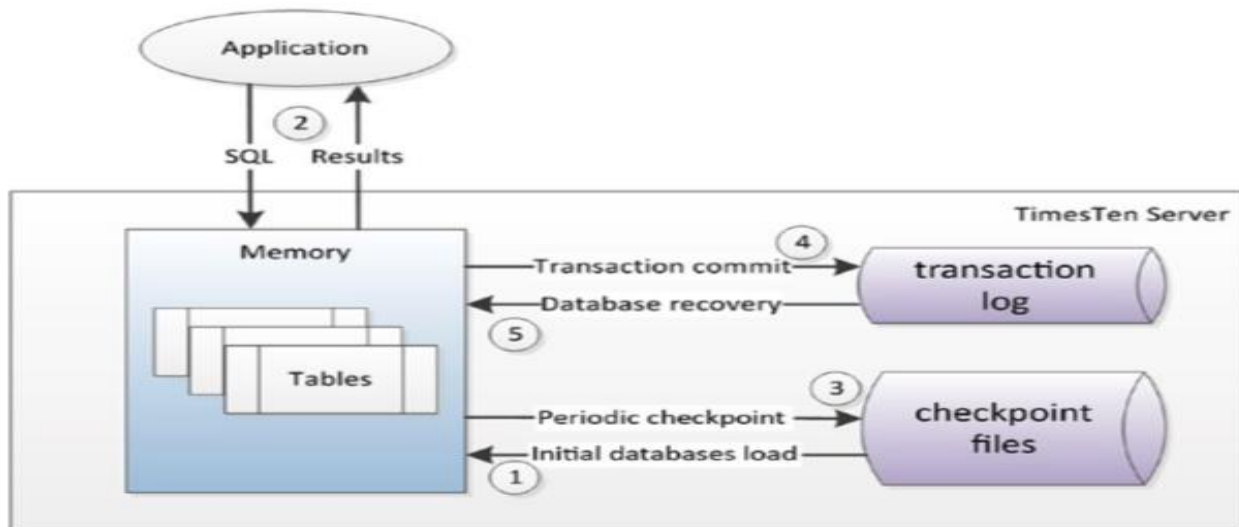


Figure 7-5. TimesTen Architecture

Redis (2009 ~ 2013 VMWare)

- Key-value store architecture

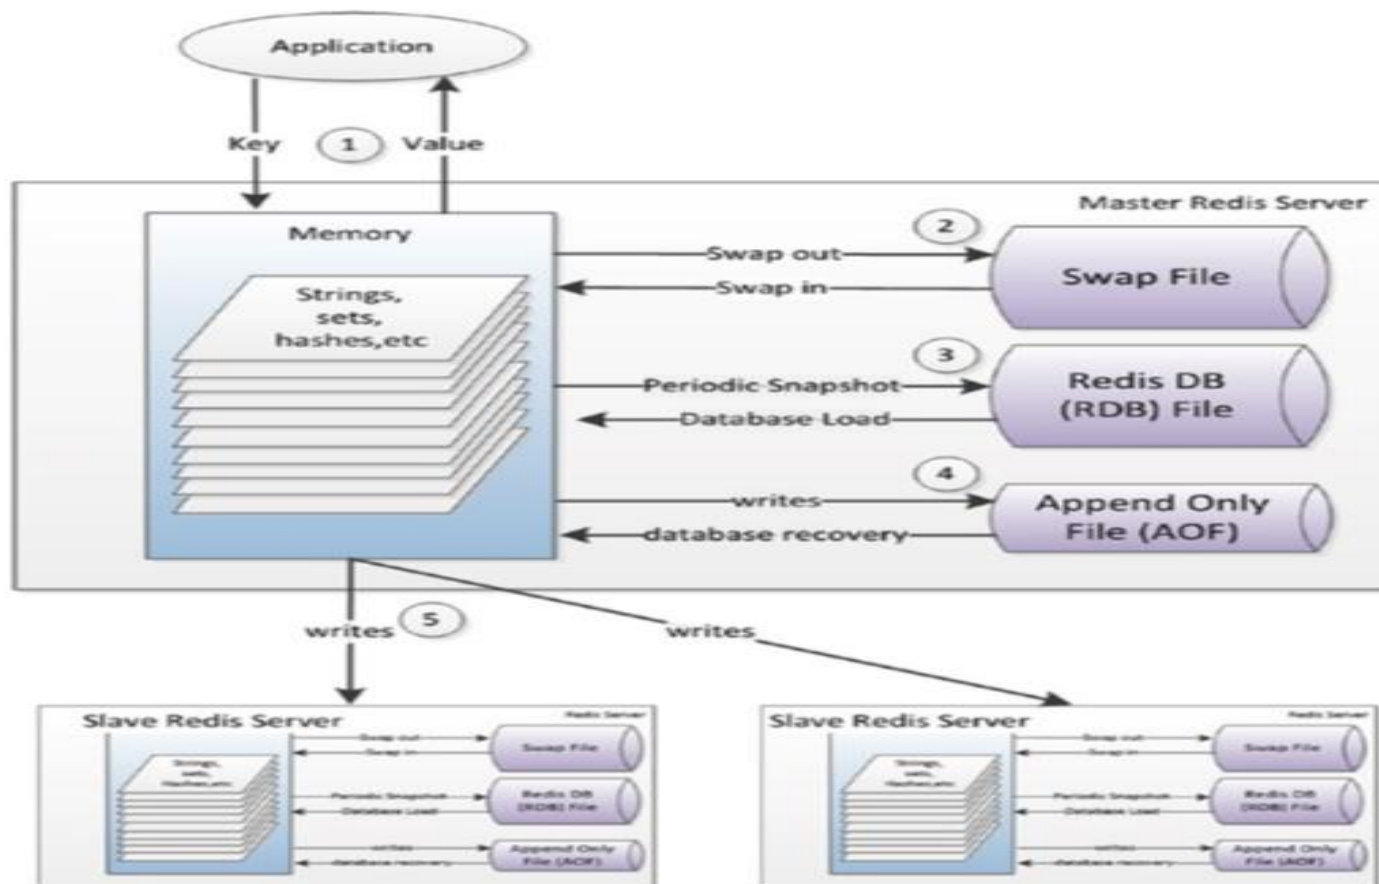


Figure 7-6. Redis architecture

SAP HANA

- In-memory database for Business Intelligence (BI)
- Capable of supporting OLTP workloads
- Column and Row Storage

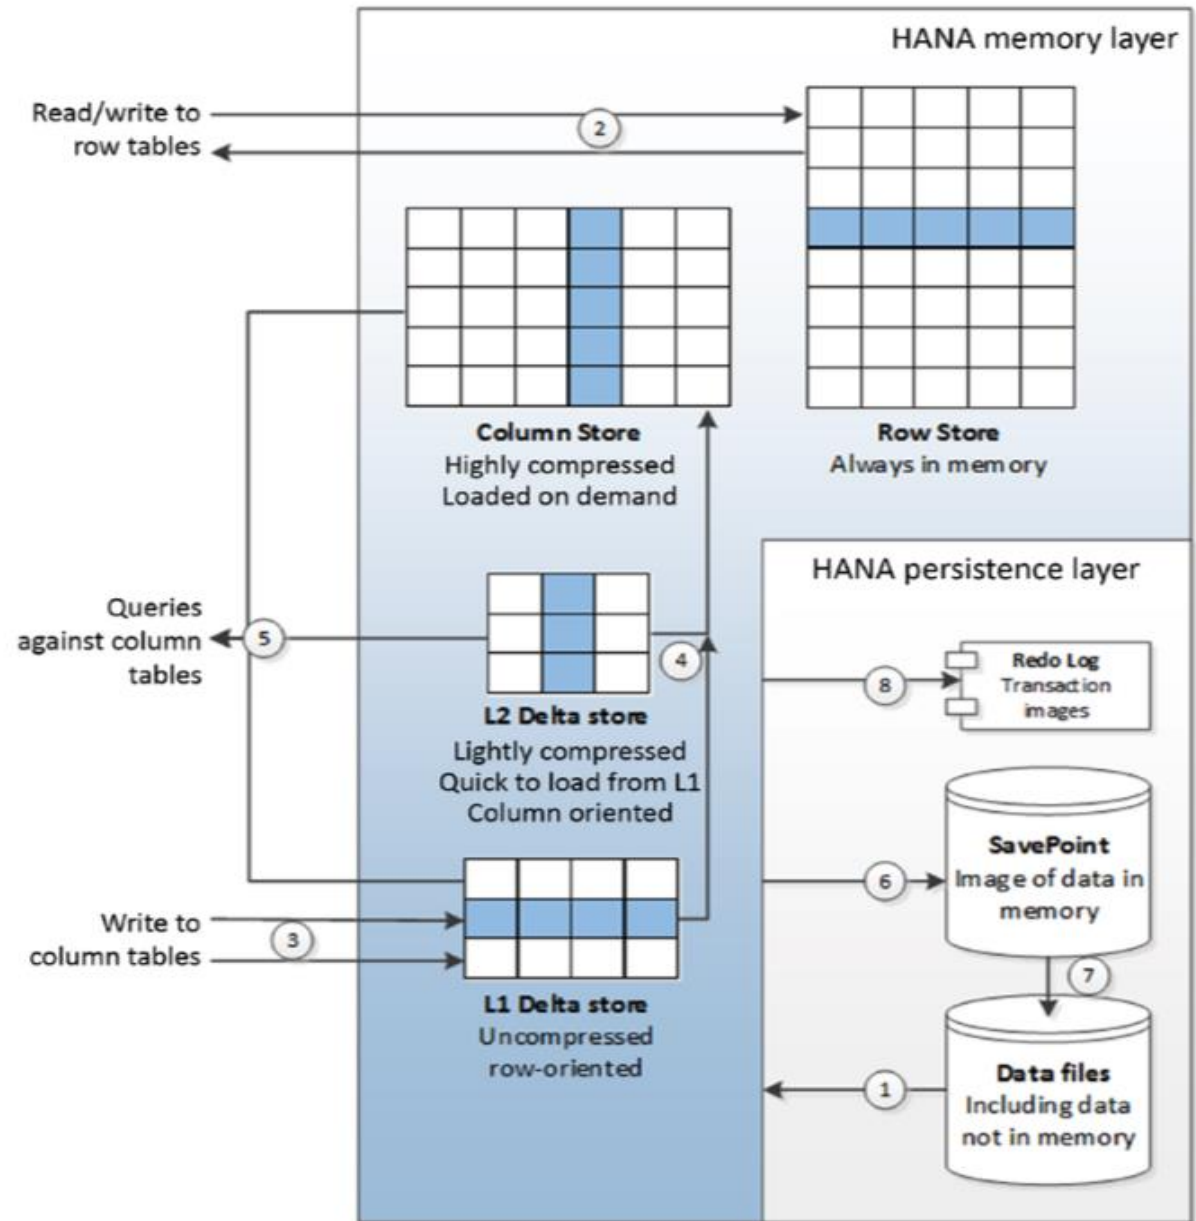


Figure 7-7. SAP HANA architecture

More in-memory databases

- VoltDB
 - Follows *H-store* design
 - A pure in-memory solution
 - Supports the ACID transactional model
 - Persistence through replication across multiple machines
- Oracle 12c “in-Memory Database”
 - In-memory column store to supplement its disk-based row store

Berkeley Analytics Data Stack (BDAS)

- **Spark**

- In-memory, distributed, fault-tolerant processing framework.
- Implemented in Scala
- Higher-level abstractions than MapReduce
- Excels at tasks that cause bottlenecks on disk IO in MapReduce
 - Good for tasks that iterate repeatedly over a dataset (e.g., machine-learning workloads)

- **Mesos**

- Cluster management layer (analogous to Hadoop's YARN)
- Intended to allow multiple frameworks

- **Tachyon**

- Fault-tolerant, Hadoop-compatible, memory-centric distributed file system

Berkeley Analytics Data Stack and Spark

- In few words:
 - *Spark represents an in-memory variation on the Hadoop theme*
 - Made to cooperate with Hadoop, not to replace it!
- BDAS includes
 - *Spark streaming*
 - *Streams analytics*
 - *GraphX*
 - Graph computation based on Spark
 - *MLBase*
 - Machine-learning libraries at various levels of abstraction

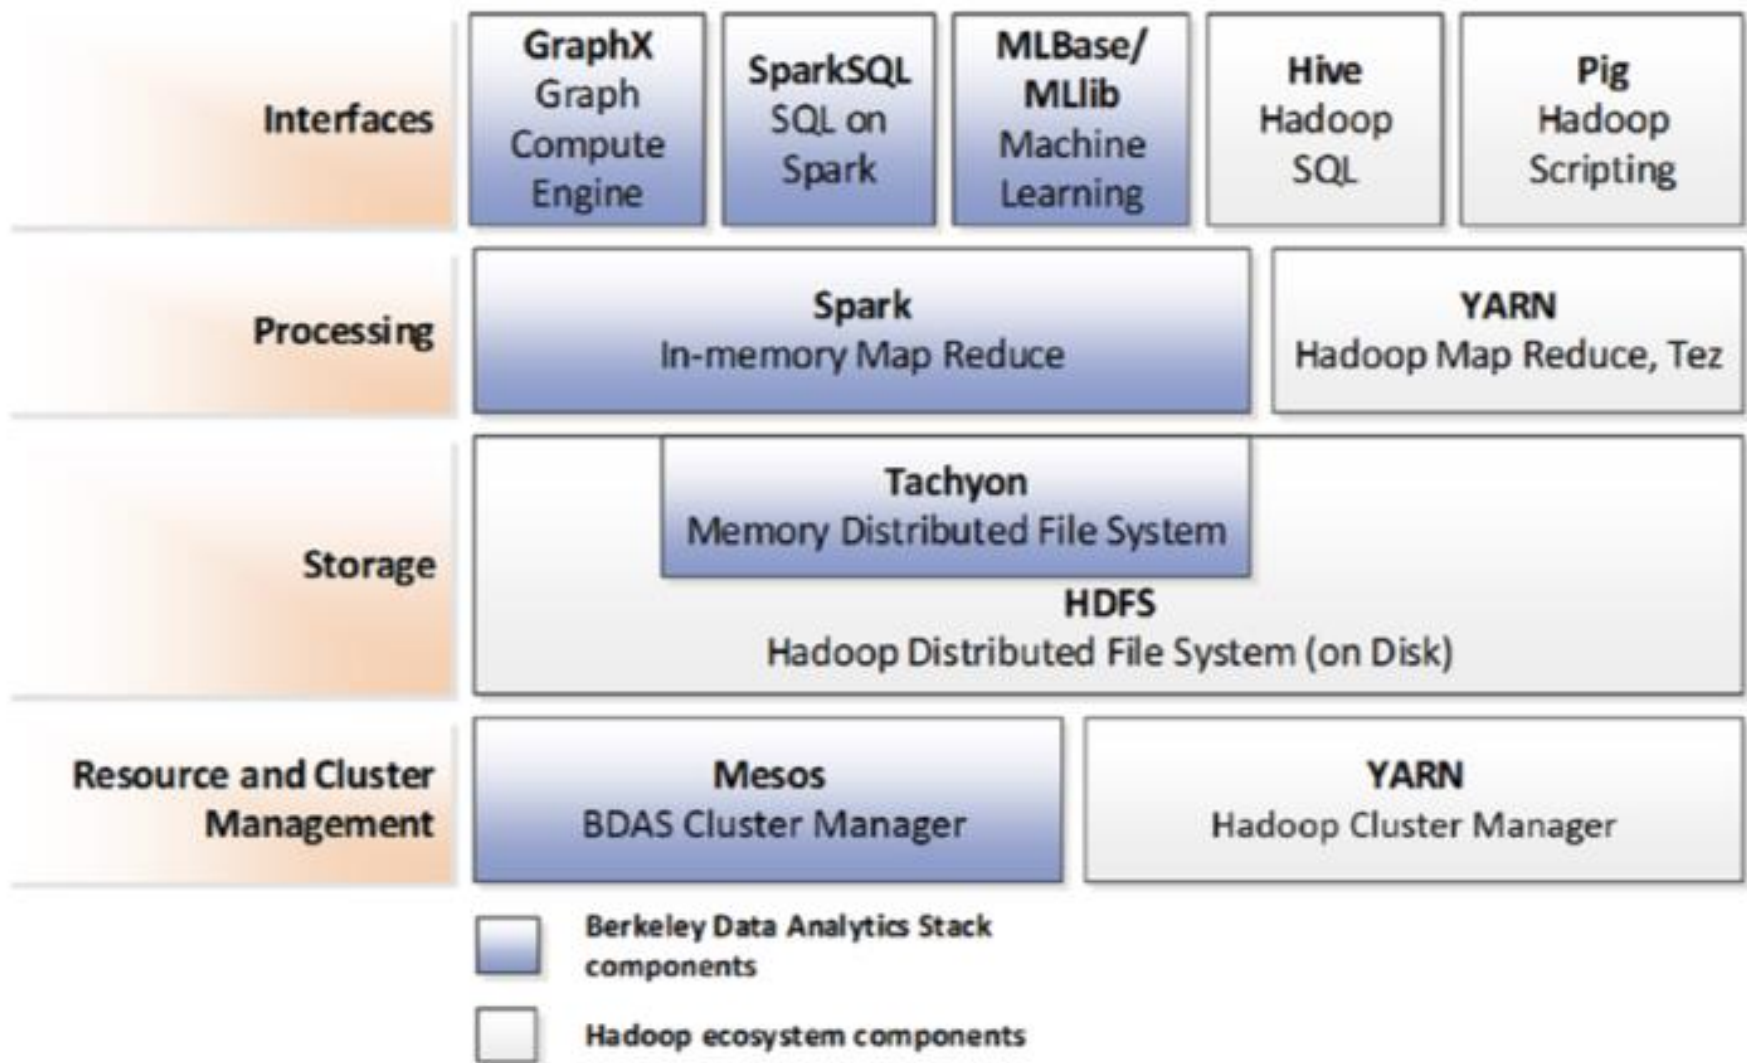


Figure 7-10. Spark, Hadoop, and the Berkeley Data Analytics Stack

Spark (1)

- Data as *resilient distributed datasets (RDD)*
 - Collections of objects
 - Simple types such as strings or any other Java/Scala object type
 - Partitioned automatically across nodes of the cluster
 - RDDs are immutable
 - Operations on RDDs return new RDDs
- Methods implemented by (*DAG*) operations
 - Directed acyclic graph
 - A more sophisticated paradigm than MapReduce
- Spark is not an OLTP-oriented system
 - No transaction log or journal

Spark (2)

- Manages data that won't fit entirely into main memory
 - May page data to disk if the data volumes exceed memory capacity
- Can read from or write to local or distributed file systems
 - It integrates with Hadoop, works with HDFS or external data
 - RDDs may be represented on disk as text files or JSON documents
- Can access data in any JDBC-compliant database
- The Spark API defines high-level methods
 - Joining, filtering, and aggregation

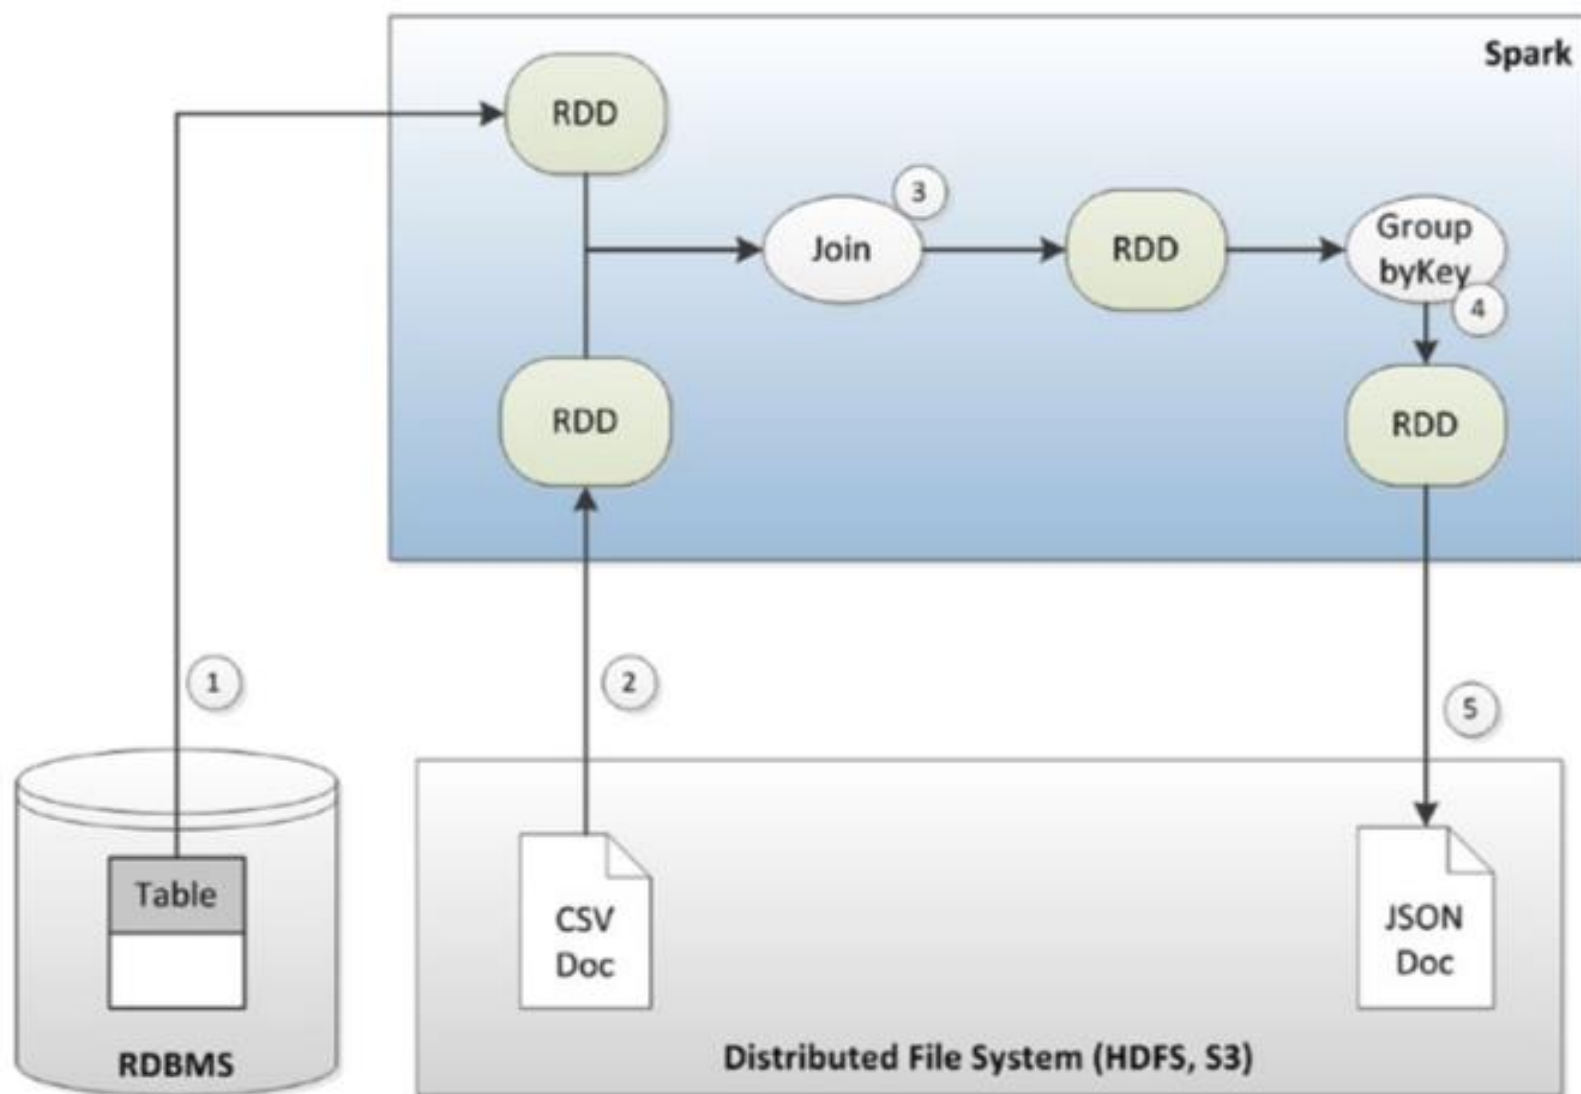


Figure 7-11. Elements of Spark processing